Please submit your solutions for this assignment on Canvas by **Friday, Nov. 22nd**. Post both **your code** and **a screenshot of its output** along with the answer to the bolded question below (about the number of weights).

We now wish to build and train a dense — that is, fully-connected — feed-forward neural net to solve the general digit classification problem. Let us agree to go ahead and use a non-linear model. This means that we need to **split out and evaluate on testing data** to gauge true performance on data unseen during training.

Please refer to Exercise 7 (on this DL@DU project page) where we define a so-called LogSoftmax model.

As discussed in class, piping the output of our hidden layers through softmax allows us to naturally interpret of the output as a *discrete probability distribution*. Since in our current project involves 10 categories, softmax looks like this:

$$\text{softmax}(x_1, x_2, \ldots, x_{10}) = \left(\frac{e^{x_1}}{\sum_{i=1}^{10} e^{x_i}}, \frac{e^{x_2}}{\sum_{i=1}^{10} e^{x_i}}, \ldots, \frac{e^{x_{10}}}{\sum_{i=1}^{10} e^{x_i}} \right).$$

However, for optimal training in practice we output through log_softmax and use NLLLoss negative log-likelihood loss as outlined here.

You may use the model on the DL@DU page (which has a single hidden layer of width 200) or any non-linear dense model you wish. As part of your solution, **compute (or approximate) and post the number of weights (i.e., trainable parameters) in your model**.

Notes:

- Since we are using a non-linear model, we need to split out test data and use it to validate our trained model.

    Suppose, as usual, that the features of our data are housed in a tensor xss in such a way that the examples are indexed by the first *axis*, and that the corresponding targets live in yss.

    In our current digit classification problem, xss is a torch.Tensor of size 5000x400, while yss is a torch.LongTensor of size 5000.

    A sensible and readable way to perform a 80/20 train/test split is:

    ```
    indices = torch.randperm(len(xss))
    xss = xss[indices]; yss = yss[indices]  # coherently randomize the data
    xss_train = xss[:4000]; yss_train = yss[:4000]
    xss_test = xss[4000:]; yss_test = yss[4000:]
    ```

- In a classification problem, you do **not** need to center or normalize the outputs yss.

- You should likely center the features xss but there is no real need to normalize them.

- You can use DUlib's train function if you wish:

    ```
    import du.lib as dulib
    ...

    model = dulib.train(
        model = model,
    ```

```
      crit = nn.NLLLoss(),
      train_data = (xss_train, yss_train),
      valid_data = (xss_test, yss_test),
      #
      #  your other training parameters
      #
  )
  ...
```

To graph, you can add the parameter graph=1.

- Feel free to use Dulib's class_accuracy function:

```
# gauge performance on training data
pct_training = dulib.class_accuracy(model, (xss_train, yss_train), show_cm=False)
print(f"Percentage correct on training data: {100*pct_training:.2f}")
```

When validating on test data, you may wish to print the confusion matrix by setting show_cm=True.